

Hands-on training session 2

Hui-Walter models for diagnostic test evaluation

Matt Denwood Giles Innocent

2020-02-18

Introduction

Overview

Date/time:

- 19th February 2020
- 16.00 - 17.00

Teachers:

- Matt Denwood (presenter)
- Giles Innocent

Recap

- Fitting models using MCMC is easy with JAGS / runjags
- But we must *never forget* to check convergence and effective sample size!
- More complex models become easy to implement
 - For example imperfect diagnostic tests
 - But remember to be realistic about what is possible with your data

Recap

- Fitting models using MCMC is easy with JAGS / runjags
- But we must *never forget* to check convergence and effective sample size!
- More complex models become easy to implement
 - For example imperfect diagnostic tests
 - But remember to be realistic about what is possible with your data
- So how do we extend these models to multiple diagnostic tests?

Session 2a: Hui-Walter models for 2 tests and 1 population

Hui-Walter Model

- A particular model formulation that was originally designed for evaluating diagnostic tests in the absence of a gold standard
- Not necessarily (or originally) Bayesian but now usually implemented using Bayesian MCMC
- But evaluating an imperfect test against another imperfect test is a bit like pulling a rabbit out of a hat
 - If we don't know the true disease status, how can we estimate sensitivity or specificity for either test?

Model Specification

```
model{
  Tally ~ dmulti(prob, TotalTests)

  # Test1- Test2-
  prob[1] <- (prev * ((1-se[1])*(1-se[2]))) + ((1-prev) *

  # Test1+ Test2-
  prob[2] <- (prev * ((se[1])*(1-se[2]))) + ((1-prev) *

  # Test1- Test2+
  prob[3] <- (prev * ((1-se[1])*(se[2]))) + ((1-prev) *

  # Test1+ Test2+
  prob[4] <- (prev * ((se[1])*(se[2]))) + ((1-prev) * (1 - (se[1] * se[2])))
```



```
prev ~ dbeta(1, 1)
se[1] ~ dbeta(1, 1)
sp[1] ~ dbeta(1, 1)
se[2] ~ dbeta(1, 1)
sp[2] ~ dbeta(1, 1)
```

```
#data# Tally, TotalTests
#monitor# prev, prob, se, sp
#inits# prev, se, sp
```

```
}
```

```
1 twoXtwo <- matrix(c(48, 12, 4, 36), ncol=2, nrow=2)
2 twoXtwo
```

```
##      [,1] [,2]
## [1,]   48   4
## [2,]   12  36
```

```
1 library('runjags')
2
3 Tally <- as.numeric(twoXtwo)
4 TotalTests <- sum(Tally)
5
6 prev <- list(chain1=0.05, chain2=0.95)
7 se <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
8 sp <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
9
10 results <- run.jags('basic_hw.bug', n.chains=2)
```

```
## Warning: You should update the rjags package to version
## (This warning is given once per R session)
```

1 results

	Lower95	Median	Upper95	SSeff	psrf
prev	0.311	0.440	0.576	3967	1
prob[1]	0.368	0.461	0.558	13300	1
prob[2]	0.074	0.133	0.204	13181	1
prob[3]	0.018	0.055	0.104	9780	1
prob[4]	0.257	0.344	0.441	13238	1
se[1]	0.821	0.932	1.000	5569	1
se[2]	0.687	0.848	1.000	3111	1
sp[1]	0.748	0.877	1.000	3179	1
sp[2]	0.861	0.948	1.000	5765	1

- Note the wide confidence intervals!

- These models need A LOT of data
 - And/or strong priors for one of the tests
- Convergence is more problematic than usual
- Be **very** careful with the order of combinations in `dmultinom`!
- Check your results carefully to ensure they make sense!

Label Switching

How to interpret a test with $Se=0\%$ and $Sp=0\%$?

Label Switching

How to interpret a test with $Se=0\%$ and $Sp=0\%$?

- The test is perfect - we are just holding it upside down. . .

Label Switching

How to interpret a test with $Se=0\%$ and $Sp=0\%$?

- The test is perfect - we are just holding it upside down...

We can force $se+sp \geq 1$:

```
1 se[1] ~ dbeta(1, 1)
2 sp[1] ~ dbeta(1, 1)T(1-se[1], )
```

Or:

```
1 se[1] ~ dbeta(1, 1)T(1-sp[1], )
2 sp[1] ~ dbeta(1, 1)
```

But not both!

This allows the test to be useless, but not worse than useless

Simulating data

Analysing simulated data is useful to check that we can recover parameter values.

```
1  se1 <- 0.9; sp1 <- 0.95; sp2 <- 0.99; se2 <- 0.8
2  prevalence <- 0.5; N <- 100
3
4  truestatus <- rbinom(N, 1, prevalence)
5  Test1 <- rbinom(N, 1, (truestatus * se1) + ((1-truestatus) *
   ↪ (1-sp1)))
6  Test2 <- rbinom(N, 1, (truestatus * se2) + ((1-truestatus) *
   ↪ (1-sp2)))
7
8  twoXtwo <- table(Test1, Test2)
9  twoXtwo
```

```
##      Test2
## Test1  0  1
##      0 38  1
```


Exercise

Modify JAGS code to force tests to be better than useless

Simulate data and recover parameters for:

- $N=10$, $N=100$, $N=1000$

Optional Exercise

Use priors for test1 taken from session 1 and compare the results

Session 2b: Hui-Walter models for 2 tests and N populations

Hui-Walter models with multiple populations

- Basically an extension of the single-population model
- Works best with multiple populations each with differing prevalences
 - These could even be subgroups of individuals within the same population if there are e.g. known risk factors for disease status
 - Remember that the focus is usually to estimate the diagnostic test parameters and not the prevalence in the different populations/subgroups!

Independent intercepts for populations

```
1  model{
2    for(p in 1:Populations){
3      Tally[1:4, p] ~ dmulti(prob[1:4, p], TotalTests[p])
4      # Test1- Test2- Pop1
5      prob[1, p] <- (prev[p] * ((1-se[1])*(1-se[2]))) +
↪    ((1-prev[p]) * ((sp[1])*(sp[2])))
6      ## etc ##
7      prev[p] ~ dbeta(1, 1)
8    }
9
10   se[1] ~ dbeta(HPSe[1,1], HPSe[1,2])T(1-sp[1], )
11   sp[1] ~ dbeta(HPSp[1,1], HPSp[1,2])
12   se[2] ~ dbeta(HPSe[2,1], HPSe[2,2])T(1-sp[2], )
13   sp[2] ~ dbeta(HPSp[2,1], HPSp[2,2])
14
15   #data# Tally, TotalTests, Populations, HPSe, HPSp
16   #monitor# prev, prob, se, sp
17   #inits# prev, se, sp
18 }
```

We would usually start with individual-level data in a dataframe:

```
1  se1 <- 0.9; sp1 <- 0.95; sp2 <- 0.99; se2 <- 0.8
2  prevalences <- c(0.1, 0.5, 0.9)
3  N <- 100
4
5  simdata <- data.frame(Population = sample(seq_along(prevalences),
   ↪  N, replace=TRUE))
6  simdata$probability <- prevalences[simdata$Population]
7  simdata$truestatus <- rbinom(N, 1, simdata$probability)
8  simdata$Test1 <- rbinom(N, 1, (simdata$truestatus * se1) +
   ↪  ((1-simdata$truestatus) * (1-sp1)))
9  simdata$Test2 <- rbinom(N, 1, (simdata$truestatus * se2) +
   ↪  ((1-simdata$truestatus) * (1-sp2)))
10
11 head(simdata)
```

##	Population	probability	truestatus	Test1	Test2
## 1	3	0.9	1	1	0
## 2	1	0.1	0	0	0
## 3	2	0.5	0	0	0
## 4	2	0.5	0	0	0
## 5	1	0.1	0	0	0
## 6	1	0.1	0	1	0

[Except that probability and truestatus would not normally be known!]

- The model code and data format for an arbitrary number of populations (and tests) can be determined automatically
- There is a function (provided in the GitHub repo) that can do this for us:

```
1 simdata$Population <- factor(simdata$Population,  
  ↪ levels=seq_along(prevalences), labels=paste0('Pop_',  
  ↪ seq_along(prevalences)))  
2  
3 source("autohuiwalter.R")  
4 auto_huiwalter(simdata[,c('Population', 'Test1', 'Test2')],  
  ↪ outfile='autohw.bug')
```

The model and data have been written to autohw.bug in the
You should check and alter priors before running the model

- This generates self-contained model/data/initial values etc:
 - [ignore covse and covsp for now]

```
## ## Auto-generated Hui-Walter model created by script ver
##
## model{
##
## ## Observation layer:
##
## # Complete observations (N=100):
## for(p in 1:Populations){
##     Tally_RR[1:4,p] ~ dmulti(prob_RR[1:4,p], N_RR[p])
##
##     prob_RR[1:4,p] <- se_prob[1:4,p] + sp_prob[1:4,p]
## }
##
##
```

And can be run directly from R:

```
1 results <- run.jags('autohw.bug')
```

```
## Note: The monitored variables 'covse12' and 'covsp12'  
## appear to be non-stochastic; they will not be  
## included in the convergence diagnostic  
## Finished running the simulation
```

```
1 results
```

```
##
```

```
## JAGS model summary statistics from 20000 samples (chains
```

```
##
```

##	Lower95	Median	Upper95	Mean	SD
## se[1]	0.81441	0.92044	0.99704	0.91299	0.050702
## se[2]	0.70145	0.84636	0.9787	0.84086	0.072665
## sp[1]	0.78389	0.88872	0.97862	0.88467	0.05079
## sp[2]	0.90831	0.97034	1	0.96371	0.028238
## prev[1]	0.0061492	0.084139	0.20415	0.093738	0.054871
## prev[2]	0.1555	0.31548	0.50268	0.3207	0.089942
## prev[3]	0.69238	0.84704	0.97197	0.83831	0.07467
## covse12	0	0	0	0	0
## covsp12	0	0	0	0	0
##					

- Modifying priors must still be done directly in the model file
- The model needs to be re-generated if the data changes
 - But remember that your modified priors will be reset
- There must be a single column for the population (as a factor), and all of the other columns (either factor, logical or numeric) are interpreted as being test results
- The function will soon be included in the runjags package
 - Feedback welcome!

Observation-level model specification

```
model{  
  
  for(i in 1:N){  
    Status[i] ~ dcat(prob[i, ])  
  
    prob[i,1] <- (prev[i] * ((1-se[1])*(1-se[2]))) +  
                 ((1-prev[i]) * ((sp[1])*(sp[2])))  
    prob[i,2] <- (prev[i] * ((se[1])*(1-se[2]))) +  
                 ((1-prev[i]) * ((1-sp[1])*(sp[2])))  
    prob[i,3] <- (prev[i] * ((1-se[1])*(se[2]))) +  
                 ((1-prev[i]) * ((sp[1])*(1-sp[2])))  
    prob[i,4] <- (prev[i] * ((se[1])*(se[2]))) +  
                 ((1-prev[i]) * ((1-sp[1])*(1-sp[2])))  
  }  
}
```

```

intercept ~ dnorm(0, 0.33)
population_effect[1] <- 0
for(p in 2:Pops){
  population_effect[p] ~ dnorm(0, 0.1)
}
se[1] ~ dbeta(1, 1)T(1-sp[1], )
sp[1] ~ dbeta(1, 1)
se[2] ~ dbeta(1, 1)T(1-sp[2], )
sp[2] ~ dbeta(1, 1)

#data# Status, N, Population, Pops
#monitor# intercept, population_effect, se, sp
#inits# intercept, population_effect, se, sp
}

```

- The main difference is the prior for prevalence in each population
- We also need to give initial values for intercept and population_effect rather than prev, and tell run.jags the data frame from which to extract the data (except N and Pops):

```
1 intercept <- list(chain1=-1, chain2=1)
2 population_effect <- list(chain1=c(NA, 1, -1), chain2=c(NA, -1,
  ↪ 1))
3 se <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
4 sp <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
5
6 simdata$Status <- with(simdata, factor(interaction(Test1, Test2),
  ↪ levels=c('0.0','1.0','0.1','1.1')))
7 N <- nrow(simdata)
8 Pops <- length(levels(simdata$Population))
9 glm_results <- run.jags('glm_hw.bug', n.chains=2, data=simdata)
```

```
## Note: The monitored variable 'population effect[1]'
```

Also like in session 1, the estimates for se/sp should be similar, although this model runs more slowly.

Note: this model could be used as the basis for adding covariates

For a handy way to generate a GLM model see
`runjags::template.jags`

- Look out for integration with `autohuiwalter` in the near (ish) future. . .

Need to be very careful with tabulating the data, or use automatically generated code

Works best when populations have very different prevalences

Exercise

Play around with the `autohuiwalter` function

Notice the model and data and initial values are in a self contained file

Ignore the `covse` and `covsp` for now

[There is no particular solution to this exercise!]

Summary

Summary

- Using JAGS / runjags allows us to work with MCMC more easily, safely and efficiently than writing our own sampling algorithms
- But we must *never forget* to check convergence and effective sample size!
- More complex models become easy to implement
 - For example imperfect diagnostic tests
- But just because a model can be defined does not mean that it will be useful for our data
 - We need to be realistic about the information available in the data, what parameters are feasible to estimate, and where we will need to use strong priors