

Hands-on training session 2

Hui-Walter models for diagnostic test evaluation

Matt Denwood Giles Innocent

2020-02-19

Introduction

Date/time:

- 19th February 2020
- 16.00 - 17.00

Teachers:

- Matt Denwood (presenter)
- Giles Innocent

Recap

- Fitting models using MCMC is easy with JAGS / runjags
- But we must **never forget** to check convergence and effective sample size!
- More complex models become easy to implement
 - For example imperfect diagnostic tests
 - But remember to be realistic about what is possible with your data

Recap

- Fitting models using MCMC is easy with JAGS / runjags
- But we must **never forget** to check convergence and effective sample size!
- More complex models become easy to implement
 - For example imperfect diagnostic tests
 - But remember to be realistic about what is possible with your data
- So how do we extend these models to multiple diagnostic tests?

Session 2a: Hui-Walter models for 2 tests and 1 population

Hui-Walter Model

- A particular model formulation that was originally designed for evaluating diagnostic tests in the absence of a gold standard
- Not necessarily (or originally) Bayesian but now usually implemented using Bayesian MCMC
- But evaluating an imperfect test against another imperfect test is a bit like pulling a rabbit out of a hat
 - If we don't know the true disease status, how can we estimate sensitivity or specificity for either test?

Model Specification

```
1  model{
2    Tally ~ dmulti(prob, TotalTests)
3
4    # Test1- Test2-
5    prob[1] <- (prev * ((1-se[1])*(1-se[2]))) + ((1-prev) *
6      ↪ ((sp[1])*(sp[2])))
7
8    # Test1+ Test2-
9    prob[2] <- (prev * ((se[1])*(1-se[2]))) + ((1-prev) *
10     ↪ ((1-sp[1])*(sp[2])))
11
12    # Test1- Test2+
13    prob[3] <- (prev * ((1-se[1])*(se[2]))) + ((1-prev) *
14     ↪ ((sp[1])*(1-sp[2])))
```



```

1
2 # Test1+ Test2+
3 prob[4] <- (prev * ((se[1])*(se[2]))) + ((1-prev) *
   ↪ ((1-sp[1])*(1-sp[2])))
4
5 prev ~ dbeta(1, 1)
6 se[1] ~ dbeta(1, 1)
7 sp[1] ~ dbeta(1, 1)
8 se[2] ~ dbeta(1, 1)
9 sp[2] ~ dbeta(1, 1)
10
11 #data# Tally, TotalTests
12 #monitor# prev, prob, se, sp
13 #inits# prev, se, sp
14 }

```

```

1  twoXtwo <- matrix(c(48, 12, 4, 36), ncol=2, nrow=2)
2  twoXtwo

1  ##      [,1] [,2]
2  ## [1,]   48   4
3  ## [2,]   12  36

1  library('runjags')
2
3  Tally <- as.numeric(twoXtwo)
4  TotalTests <- sum(Tally)
5
6  prev <- list(chain1=0.05, chain2=0.95)
7  se <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
8  sp <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
9
10 results <- run.jags('basic_hw.bug', n.chains=2)

```

[Remember to check convergence and effective sample size!]

1 results

	Lower95	Median	Upper95	SSeff	psrf
prev	0.309	0.444	0.572	3858	1.000
prob[1]	0.369	0.461	0.562	13053	1.000
prob[2]	0.074	0.133	0.204	14426	1.000
prob[3]	0.019	0.055	0.104	9757	1.000
prob[4]	0.254	0.344	0.434	13271	1.000
se[1]	0.823	0.931	1.000	5688	1.000
se[2]	0.684	0.842	1.000	3279	1.001
sp[1]	0.747	0.882	1.000	3348	1.001
sp[2]	0.863	0.948	1.000	5640	1.001

- Note the wide confidence intervals!

- These models need A LOT of data
 - And/or strong priors for one of the tests
- Convergence is more problematic than usual
- Be **very** careful with the order of combinations in `dmultinom`!
- Check your results carefully to ensure they make sense!

Label Switching

How to interpret a test with $Se=0\%$ and $Sp=0\%$?

Label Switching

How to interpret a test with $Se=0\%$ and $Sp=0\%$?

- The test is perfect - we are just holding it upside down. . .

Label Switching

How to interpret a test with $Se=0\%$ and $Sp=0\%$?

- The test is perfect - we are just holding it upside down...

We can force $se+sp \geq 1$:

```
1 se[1] ~ dbeta(1, 1)
2 sp[1] ~ dbeta(1, 1)T(1-se[1], )
```

Or:

```
1 se[1] ~ dbeta(1, 1)T(1-sp[1], )
2 sp[1] ~ dbeta(1, 1)
```

But not both!

This allows the test to be useless, but not worse than useless

Simulating data

Analysing simulated data is useful to check that we can recover parameter values.

```
1  se1 <- 0.9; sp1 <- 0.95;
2  se2 <- 0.8; sp2 <- 0.99
3  prevalence <- 0.5; N <- 100
4
5  truestatus <- rbinom(N, 1, prevalence)
6  Test1 <- rbinom(N, 1, (truestatus * se1) + ((1-truestatus) *
   ↪ (1-sp1)))
7  Test2 <- rbinom(N, 1, (truestatus * se2) + ((1-truestatus) *
   ↪ (1-sp2)))
8
9  twoXtwo <- table(Test1, Test2)
10 Tally <- as.numeric(twoXtwo)
```

Can we recover these parameter values?

Exercise

Modify the code in the Hui Walter model to force tests to be no worse than useless

Simulate data and recover parameters for:

- $N=10$
- $N=100$
- $N=1000$

Optional Exercise

Compare results with the following priors for test 1:

- Sensitivity = 0.9 (95% CI: 0.85 - 0.95)
- Specificity = 0.95 (95%CI: 0.92-0.97)

[These are the same as in session 1]

Session 2b: Hui-Walter models for 2 tests and N populations

Hui-Walter models with multiple populations

- Basically an extension of the single-population model
- Works best with multiple populations each with differing prevalences
 - These could even be subgroups of individuals within the same population if there are e.g. known risk factors for disease status
 - Remember that the focus is usually to estimate the diagnostic test parameters and not the prevalence in the different populations/subgroups!

Independent intercepts for populations

```
1  model{
2    for(p in 1:Populations){
3      Tally[1:4, p] ~ dmulti(prob[1:4, p], TotalTests[p])
4      # Test1- Test2- Pop1
5      prob[1, p] <- (prev[p] * ((1-se[1])*(1-se[2]))) +
↪    ((1-prev[p]) * ((sp[1])*(sp[2])))
6      ## etc ##
7      prev[p] ~ dbeta(1, 1)
8    }
9
10   se[1] ~ dbeta(HPSe[1,1], HPSe[1,2])T(1-sp[1], )
11   sp[1] ~ dbeta(HPSp[1,1], HPSp[1,2])
12   se[2] ~ dbeta(HPSe[2,1], HPSe[2,2])T(1-sp[2], )
13   sp[2] ~ dbeta(HPSp[2,1], HPSp[2,2])
14
15   #data# Tally, TotalTests, Populations, HPSe, HPSp
16   #monitor# prev, prob, se, sp
17   #inits# prev, se, sp
18 }
```

We would usually start with individual-level data in a dataframe:

```
1  se1 <- 0.9; sp1 <- 0.95; sp2 <- 0.99; se2 <- 0.8
2  prevalences <- c(0.1, 0.5, 0.9)
3  N <- 100
4
5  simdata <- data.frame(Population = sample(seq_along(prevalences),
   ↪ N, replace=TRUE))
6  simdata$probability <- prevalences[simdata$Population]
7  simdata$truestatus <- rbinom(N, 1, simdata$probability)
8  simdata$Test1 <- rbinom(N, 1, (simdata$truestatus * se1) +
   ↪ ((1-simdata$truestatus) * (1-sp1)))
9  simdata$Test2 <- rbinom(N, 1, (simdata$truestatus * se2) +
   ↪ ((1-simdata$truestatus) * (1-sp2)))
```

```
1 head(simdata)
```

```
1  ##      Population probability truestatus Test1 Test2
2  ## 1           2           0.5           0      0      0
3  ## 2           3           0.9           1      1      1
4  ## 3           1           0.1           0      0      0
5  ## 4           2           0.5           1      1      1
6  ## 5           3           0.9           1      0      1
7  ## 6           2           0.5           0      0      0
```

[Except that probability and truestatus would not normally be known!]

The model code and data format for an arbitrary number of populations (and tests) can be determined automatically

There is a function (provided in the GitHub repo) that can do this for us:

```
1  simdata$Population <- factor(simdata$Population,  
  ↪  levels=seq_along(prevalences), labels=paste0('Pop_',  
  ↪  seq_along(prevalences)))  
2  
3  source("autohuiwalter.R")  
4  auto_huiwalter(simdata[,c('Population', 'Test1', 'Test2')],  
  ↪  outfile='autohw.bug')
```


This generates self-contained model/data/initial values etc (ignore covse and covsp for now):

```
1  model{
2
3      ## Observation layer:
4
5      # Complete observations (N=100):
6      for(p in 1:Populations){
7          Tally_RR[1:4,p] ~ dmulti(prob_RR[1:4,p], N_RR[p])
8
9          prob_RR[1:4,p] <- se_prob[1:4,p] + sp_prob[1:4,p]
10     }
11
12
13     ## Observation probabilities:
14
15     for(p in 1:Populations){
16
17         # Probability of observing Test1- Test2- from a true
18         ↪ positive::
19         se_prob[1,p] <- prev[p] * ((1-se[1])*(1-se[2]) +covse12)
20         # Probability of observing Test1- Test2- from a true
```

And can be run directly from R:

```
1 results <- run.jags('autohw.bug')
2 results
```

	Lower95	Median	Upper95	SSeff	psrf
se[1]	0.752	0.865	0.961	9385	1.001
se[2]	0.772	0.893	1.000	6822	1.000
sp[1]	0.758	0.869	0.964	8132	1.000
sp[2]	0.928	0.983	1.000	4156	1.000
prev[1]	0.000	0.057	0.159	6421	1.000
prev[2]	0.261	0.428	0.593	11018	1.000
prev[3]	0.692	0.850	0.973	7781	1.000
covse12	0.000	0.000	0.000	NA	NA
covsp12	0.000	0.000	0.000	NA	NA

- Modifying priors must still be done directly in the model file
- The model needs to be re-generated if the data changes
 - But remember that your modified priors will be reset
- There must be a single column for the population (as a factor), and all of the other columns (either factor, logical or numeric) are interpreted as being test results
- The function will soon be included in the runjags package
 - Feedback welcome!

Observation-level model specification

```
1  model{
2
3    for(i in 1:N){
4      Status[i] ~ dcat(prob[i, ])
5
6      prob[i,1] <- (prev[i] * ((1-se[1])*(1-se[2]))) +
7                  ((1-prev[i]) * ((sp[1])*(sp[2])))
8      prob[i,2] <- (prev[i] * ((se[1])*(1-se[2]))) +
9                  ((1-prev[i]) * ((1-sp[1])*(sp[2])))
10     prob[i,3] <- (prev[i] * ((1-se[1])*(se[2]))) +
11                 ((1-prev[i]) * ((sp[1])*(1-sp[2])))
12     prob[i,4] <- (prev[i] * ((se[1])*(se[2]))) +
13                 ((1-prev[i]) * ((1-sp[1])*(1-sp[2])))
14
15     logit(prev[i]) <- intercept +
16       ↪ population_effect[Population[i]]
17   }
```

```

1
2 intercept ~ dnorm(0, 0.33)
3 population_effect[1] <- 0
4 for(p in 2:Pops){
5     population_effect[p] ~ dnorm(0, 0.1)
6 }
7 se[1] ~ dbeta(1, 1)T(1-sp[1], )
8 sp[1] ~ dbeta(1, 1)
9 se[2] ~ dbeta(1, 1)T(1-sp[2], )
10 sp[2] ~ dbeta(1, 1)
11
12 #data# Status, N, Population, Pops
13 #monitor# intercept, population_effect, se, sp
14 #inits# intercept, population_effect, se, sp
15 }

```

- The main difference is the prior for prevalence in each population
- We also need to give initial values for intercept and population_effect rather than prev, and tell run.jags the data frame from which to extract the data (except N and Pops):

```
1 intercept <- list(chain1=-1, chain2=1)
2 population_effect <- list(chain1=c(NA, 1, -1), chain2=c(NA, -1,
  ↪ 1))
3 se <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
4 sp <- list(chain1=c(0.5,0.99), chain2=c(0.99,0.5))
5
6 simdata$Status <- with(simdata, factor(interaction(Test1, Test2),
  ↪ levels=c('0.0','1.0','0.1','1.1')))
7 N <- nrow(simdata)
8 Pops <- length(levels(simdata$Population))
9 glm_results <- run.jags('glm_hw.bug', n.chains=2, data=simdata)
```

Also like in session 1, the estimates for se/sp should be similar, although this model runs more slowly.

Note: this model could be used as the basis for adding covariates

For a handy way to generate a GLM model see
`runjags::template.jags`

- Look out for integration with `autohuiwalter` in the near (ish) future. . .

Exercise

Play around with the `autohuiwalter` function

Notice the model and data and initial values are in a self contained file

Ignore the `covse` and `covsp` for now

[There is no particular solution to this exercise!]

Summary

- Estimating sensitivity and specificity is like pulling a rabbit out of a hat
- Multiple populations helps **a lot**
- Strong priors for one of the tests helps even more!
- Make sure you tabulate the data correctly . . . or use the automated model generator!