

# Hands-on training session 3

Hui-Walter models with more than two diagnostic tests

---

Matt Denwood   Giles Innocent   Sonja Hartnack

2020-02-19

# Introduction

---

# Overview

Date/time:

- 20th February 2020
- 14.00 - 15.30

Teachers:

- Matt Denwood (presenter)
- Giles Innocent
- Sonja Hartnack

# Recap

- JAGS / runjags is the easy way to work with complex models
  - But we *still have to* check convergence and effective sample size!
- Estimating sensitivity and specificity is like pulling a rabbit out of a hat
  - Multiple populations helps **a lot**
  - Strong priors for one of the tests helps even more!

## Recap

- JAGS / runjags is the easy way to work with complex models
  - But we *still have to* check convergence and effective sample size!
- Estimating sensitivity and specificity is like pulling a rabbit out of a hat
  - Multiple populations helps **a lot**
  - Strong priors for one of the tests helps even more!
- But what if the tests are not independent of each other?

## **Session 3a: Hui-Walter models for multiple conditionally independent tests**

---

# What exactly is our latent class?

- What do we mean by “conditionally independent?”

# What exactly is our latent class?

- What do we mean by “conditionally independent?”
- Example: we have three antibody tests
  - The latent status is actually ‘producing antibodies’ not ‘diseased’



# What exactly is our latent class?

- What do we mean by “conditionally independent?”
- Example: we have three antibody tests
  - The latent status is actually ‘producing antibodies’ not ‘diseased’
- Example: antibody vs egg count tests for liver fluke
  - Does the latent state include migrating juvenile fluke?

# What exactly is our latent class?

- What do we mean by “conditionally independent?”
- Example: we have three antibody tests
  - The latent status is actually ‘producing antibodies’ not ‘diseased’
- Example: antibody vs egg count tests for liver fluke
  - Does the latent state include migrating juvenile fluke?
- We’re actually pulling **something** out of a hat, and deciding to call it a rabbit

# Simulating data

Simulating data using an arbitrary number of independent tests is quite straightforward.

```
1  # Parameter values to simulate:
2  N <- 200
3  se1 <- 0.8
4  sp1 <- 0.95
5  se2 <- 0.9
6  sp2 <- 0.99
7  se3 <- 0.95
8  sp3 <- 0.95
9
10 Populations <- 2
11 prevalence <- c(0.25,0.75)
12 Group <- sample(1:Populations, N, replace=TRUE)
```

```

1  # Ensure replicable data:
2  set.seed(2020-02-18)
3
4  # Simulate the true latent state (which is unobserved in real
   ↪ life):
5  true <- rbinom(N, 1, prevalence[Group])
6  # Simulate test results for test 1:
7  test1 <- rbinom(N, 1, se1*true + (1-sp1)*(1-true))
8  # Simulate test results for test 2:
9  test2 <- rbinom(N, 1, se2*true + (1-sp2)*(1-true))
10 # Simulate test results for test 3:
11 test3 <- rbinom(N, 1, se3*true + (1-sp3)*(1-true))
12
13 simdata <- data.frame(Population=factor(Group), Test1=test1,
   ↪ Test2=test2, Test3=test3)

```

## Model specification

- Like for two tests, except it is now a  $2 \times 2 \times 2$  table
  - If calculating this manually, take **extreme** care with multinomial tabulation

# Model specification

- Like for two tests, except it is now a 2x2x2 table
  - If calculating this manually, take **extreme** care with multinomial tabulation
- Or use `autohuiwalter`
  - This will also handle missing data in one or more test results

```
1 source("autohuiwalter.R")
2 auto_huiwalter(simdata[,c('Population', 'Test1', 'Test2', 'Test3')],
  ↪ outfile='auto3thw.bug')
```

```

1  Tally_RRR[1:8,p] ~ dmulti(prob_RRR[1:8,p], N_RRR[p])
2  prob_RRR[1:8,p] <- se_prob[1:8,p] + sp_prob[1:8,p]
3
4  # Probability of observing Test1- Test2- Test3- from a true
   ↪ positive::
5  se_prob[1,p] <- prev[p] * ((1-se[1])*(1-se[2])*(1-se[3]) +covse12
   ↪ +covse13 +covse23)
6  # Probability of observing Test1- Test2- Test3- from a true
   ↪ negative::
7  sp_prob[1,p] <- (1-prev[p]) * (sp[1]*sp[2]*sp[3] +covsp12
   ↪ +covsp13 +covsp23)
8
9  . . .
10
11 # Probability of observing Test1+ Test2+ Test3+ from a true
   ↪ positive::
12 se_prob[8,p] <- prev[p] * (se[1]*se[2]*se[3] +covse12 +covse13
   ↪ +covse23)
13 # Probability of observing Test1+ Test2+ Test3+ from a true
   ↪ negative::
14 sp_prob[8,p] <- (1-prev[p]) * ((1-sp[1])*(1-sp[2])*(1-sp[3])
   ↪ +covsp12 +covsp13 +covsp23)

```

## Alternative model specification

We might want to explicitly model the latent state:

```
1  model{
2
3    for(i in 1:N){
4      truestatus[i] ~ dbern(prev[Population[i]])
5
6      Status[i] ~ dcat(prob[1:8, i])
7      prob[1:8,i] <- se_prob[1:8,i] + sp_prob[1:8,i]
8
9      se_prob[1,i] <- truestatus[i] *
10         ↪ ((1-se[1])*(1-se[2])*(1-se[3]))
11      sp_prob[1,i] <- (1-truestatus[i]) * (sp[1]*sp[2]*sp[3])
12
13      se_prob[2,i] <- truestatus[i] *
14         ↪ (se[1]*(1-se[2])*(1-se[3]))
15      sp_prob[2,i] <- (1-truestatus[i]) *
16         ↪ ((1-sp[1])*sp[2]*sp[3])
```



```

1      se_prob[3,i] <- truestatus[i] *
2      ↪ ((1-se[1])*se[2]*(1-se[3]))
3      sp_prob[3,i] <- (1-truestatus[i]) *
4      ↪ (sp[1]*(1-sp[2])*sp[3])
5
6      se_prob[4,i] <- truestatus[i] * (se[1]*se[2]*(1-se[3]))
7      sp_prob[4,i] <- (1-truestatus[i]) *
8      ↪ ((1-sp[1])*(1-sp[2])*sp[3])
9
10     se_prob[5,i] <- truestatus[i] *
11     ↪ ((1-se[1])*(1-se[2])*se[3])
12     sp_prob[5,i] <- (1-truestatus[i]) *
13     ↪ (sp[1]*sp[2]*(1-sp[3]))
14
15     se_prob[6,i] <- truestatus[i] * (se[1]*(1-se[2])*se[3])
16     sp_prob[6,i] <- (1-truestatus[i]) *
17     ↪ ((1-sp[1])*sp[2]*(1-sp[3]))
18
19     se_prob[7,i] <- truestatus[i] * ((1-se[1])*se[2]*se[3])
20     sp_prob[7,i] <- (1-truestatus[i]) *
21     ↪ (sp[1]*(1-sp[2])*(1-sp[3]))

```

```

1
2     se_prob[8,i] <- truestatus[i] * (se[1]*se[2]*se[3])
3     sp_prob[8,i] <- (1-truestatus[i]) *
4     ↪     ((1-sp[1])*(1-sp[2])*(1-sp[3]))
5
6     }
7
8     prev[1] ~ dbeta(1,1)
9     prev[2] ~ dbeta(1,1)
10
11    se[1] ~ dbeta(1, 1)T(1-sp[1], )
12    sp[1] ~ dbeta(1, 1)
13    se[2] ~ dbeta(1, 1)T(1-sp[2], )
14    sp[2] ~ dbeta(1, 1)
15    se[3] ~ dbeta(1, 1)T(1-sp[3], )
16    sp[3] ~ dbeta(1, 1)
17
18    #data# Status, N, Population
19    #monitor# prev, se, sp, truestatus[1:5]
20    #inits# prev, se, sp
21  }

```

```

1 Population <- simdata$Population
2 Status <- with(simdata, factor(interaction(Test1, Test2, Test3),
  ↳ levels=c('0.0.0','1.0.0','0.1.0','0.0.1','1.1.0','1.0.1','0.1.1','1
3
4 prev <- list(chain1=c(0.05,0.95), chain2=c(0.95,0.05))
5 se <- list(chain1=c(0.5,0.75,0.99), chain2=c(0.99,0.5,0.75))
6 sp <- list(chain1=c(0.5,0.75,0.99), chain2=c(0.99,0.5,0.75))

```

```

1 results <- run.jags('glm_hw3t.bug', n.chains=2)

```

```

1 results

```

```

1 ##
2 ## JAGS model summary statistics from 20000 samples (chains = 2;
  ↳ adapt+burnin = 5000):

```

```

3 ##

```

```

4 ##           Lower95  Median Upper95    Mean      SD
5 ## prev[1]      0.21422 0.30021 0.39187 0.30141 0.045606
6 ## prev[2]      0.65574 0.75334 0.84256 0.75111 0.048021
7 ## se[1]         0.75699 0.83684 0.91043 0.83492 0.039756
8 ## se[2]         0.77069 0.8477 0.92018 0.84555 0.038603

```

|               | Lower95 | Median | Upper95 | SSeff | psrf  |
|---------------|---------|--------|---------|-------|-------|
| prev[1]       | 0.214   | 0.300  | 0.392   | 13448 | 1.000 |
| prev[2]       | 0.656   | 0.753  | 0.843   | 11045 | 1.000 |
| se[1]         | 0.757   | 0.837  | 0.910   | 7202  | 1.000 |
| se[2]         | 0.771   | 0.848  | 0.920   | 6576  | 1.000 |
| se[3]         | 0.874   | 0.933  | 0.986   | 5564  | 1.001 |
| sp[1]         | 0.883   | 0.939  | 0.985   | 6686  | 1.000 |
| sp[2]         | 0.923   | 0.971  | 1.000   | 3757  | 1.001 |
| sp[3]         | 0.915   | 0.969  | 1.000   | 3057  | 1.000 |
| truestatus[1] | 0.000   | 0.000  | 0.000   | 20000 | 1.002 |
| truestatus[2] | 1.000   | 1.000  | 1.000   | 20000 | 1.106 |
| truestatus[3] | 1.000   | 1.000  | 1.000   | 10000 | 1.291 |
| truestatus[4] | 1.000   | 1.000  | 1.000   | 10000 | 1.291 |
| truestatus[5] | 1.000   | 1.000  | 1.000   | 20000 | 1.010 |

But this is inefficient

- Time taken is 1.6 minutes rather than a few seconds
- And the barely stochastic nature of some truestatus estimates triggers false convergence warnings
- And there is no way to distinguish individuals within the same boxes anyway, as they have the same data!

But this is inefficient

- Time taken is 1.6 minutes rather than a few seconds
- And the barely stochastic nature of some `truestatus` estimates triggers false convergence warnings
- And there is no way to distinguish individuals within the same boxes anyway, as they have the same data!

It is much better to use the estimated `se/sp/prev` to post-calculate these `truestatus` probabilities

- This can be useful for post-hoc ROC

## Exercise

Simulate data from 3 tests and analyse using the `autohuiwalter` function

Do the estimates of  $Se/Sp$  correspond to the simulation parameters?

Make some data missing for one or more tests and re-generate the model

Can you see what has changed in the code?

## Optional Exercise

Modify the simulation code to introduce an antibody response step between the true status and the test results (see below in the HTML file for example R code).

Simulate data from three antibody tests including the antibody response step

Does the sensitivity / specificity estimated by the model recover the true prevalence parameter?



## **Session 3b: Hui-Walter models for multiple tests with conditional dependence**

---

## Branching of processes leading to test results

- Sometimes we have multiple tests that are detecting a similar thing
  - For example: two antibody tests and one antigen test
  - The antibody tests will be correlated

## Branching of processes leading to test results

- Sometimes we have multiple tests that are detecting a similar thing
  - For example: two antibody tests and one antigen test
  - The antibody tests will be correlated
- Or even three antibody tests where two are primed to detect the same thing, and one has a different target!
  - In this case all three tests are correlated, but two are more strongly correlated

# Simulating data

It helps to consider the data simulation as a biological process.

```
1  # Parameter values to simulate:
2  N <- 200
3  se1 <- 0.8; sp1 <- 0.95
4  se2 <- 0.9; sp2 <- 0.99
5  se3 <- 0.95; sp3 <- 0.95
6
7  Populations <- 2
8  prevalence <- c(0.25,0.75)
9  Group <- rep(1:Populations, each=N)
10
11 # Ensure replicable data:
12 set.seed(2017-11-21)
13
14 # The probability of an antibody response given disease:
15 abse <- 0.8
16 # The probability of no antibody response given no disease:
17 absp <- 1 - 0.2
```

```

1  # Simulate the true latent state:
2  true <- rbinom(N*Populations, 1, prevalence[Group])
3
4  # Tests 1 & 2 will be co-dependent on antibody response:
5  antibody <- rbinom(N*Populations, 1, abse*true +
   ↪ (1-absp)*(1-true))
6  # Simulate test 1 & 2 results based on this other latent state:
7  test1 <- rbinom(N*Populations, 1, se1*antibody +
   ↪ (1-sp1)*(1-antibody))
8  test2 <- rbinom(N*Populations, 1, se2*antibody +
   ↪ (1-sp2)*(1-antibody))
9
10 # Simulate test results for the independent test 3:
11 test3 <- rbinom(N*Populations, 1, se3*true + (1-sp3)*(1-true))
12
13 ind3tests <- data.frame(Population=Group, Test1=test1,
   ↪ Test2=test2, Test3=test3)

```

```
1  # The overall sensitivity of the correlated tests is:
2  abse*se1 + (1-abse)*(1-sp1)
```

```
1  ## [1] 0.65
```

```
1  abse*se2 + (1-abse)*(1-sp2)
```

```
1  ## [1] 0.722
```

```
1  # The overall specificity of the correlated tests is:
2  absp*sp1 + (1-absp)*(1-se1)
```

```
1  ## [1] 0.8
```

```
1  absp*sp2 + (1-absp)*(1-se2)
```

```
1  ## [1] 0.812
```

```
1  # The overall sensitivity of the correlated tests is:
2  abse*se1 + (1-abse)*(1-sp1)
```

```
1  ## [1] 0.65
```

```
1  abse*se2 + (1-abse)*(1-sp2)
```

```
1  ## [1] 0.722
```

```
1  # The overall specificity of the correlated tests is:
2  absp*sp1 + (1-absp)*(1-se1)
```

```
1  ## [1] 0.8
```

```
1  absp*sp2 + (1-absp)*(1-se2)
```

```
1  ## [1] 0.812
```

We need to think carefully about what we are conditioning on when interpreting sensitivity and specificity!

# Model specification

```
1      se_prob[1,p] <- prev[p] * ((1-se[1])*(1-se[2])*(1-se[3])
↪      +covse12 +covse13 +covse23)
2      sp_prob[1,p] <- (1-prev[p]) * (sp[1]*sp[2]*sp[3] +covsp12
↪      +covsp13 +covsp23)
3
4      se_prob[2,p] <- prev[p] * (se[1]*(1-se[2])*(1-se[3]) -covse12
↪      -covse13 +covse23)
5      sp_prob[2,p] <- (1-prev[p]) * ((1-sp[1])*sp[2]*sp[3] -covsp12
↪      -covsp13 +covsp23)
6
7      ...
8
9      # Covariance in sensitivity between tests 1 and 2:
10     covse12 ~ dunif( (se[1]-1)*(1-se[2]) , min(se[1],se[2]) -
↪     se[1]*se[2] )
11
12     # Covariance in specificity between tests 1 and 2:
13     covsp12 ~ dunif( (sp[1]-1)*(1-sp[2]) , min(sp[1],sp[2]) -
↪     sp[1]*sp[2] )
```



# Generating the model

First use `autohuiwalter` to create a model file:

```
1 auto_huiwalter(ind3tests, 'auto3tihw.bug')
```

Then find the lines for the covariances that we want to activate:

```
1 # Covariance in sensitivity between Test1 and Test2 tests:
2 # covse12 ~ dunif( (se[1]-1)*(1-se[2]) , min(se[1],se[2]) -
  ↪ se[1]*se[2] ) ## if the sensitivity of these tests may be
  ↪ correlated
3 covse12 <- 0 ## if the sensitivity of these tests can be
  ↪ assumed to be independent
4 # Covariance in specificity between Test1 and Test2 tests:
5 # covsp12 ~ dunif( (sp[1]-1)*(1-sp[2]) , min(sp[1],sp[2]) -
  ↪ sp[1]*sp[2] ) ## if the specificity of these tests may be
  ↪ correlated
6 covsp12 <- 0 ## if the specificity of these tests can be
  ↪ assumed to be independent
```

And edit so it looks like:

```
1 # Covariance in sensitivity between Test1 and Test2 tests:
2 covse12 ~ dunif( (se[1]-1)*(1-se[2]) , min(se[1],se[2]) -
  ↪ se[1]*se[2] ) ## if the sensitivity of these tests may be
  ↪ correlated
3 # covse12 <- 0 ## if the sensitivity of these tests can be
  ↪ assumed to be independent
4 # Covariance in specificity between Test1 and Test2 tests:
5 covsp12 ~ dunif( (sp[1]-1)*(1-sp[2]) , min(sp[1],sp[2]) -
  ↪ sp[1]*sp[2] ) ## if the specificity of these tests may be
  ↪ correlated
6 # covsp12 <- 0 ## if the specificity of these tests can be
  ↪ assumed to be independent
```

[i.e. swap the comments around]

You will also need to uncomment out the relevant initial values for BOTH chains (on lines 117-122 and 128-133):

```
1 # "covse12" <- 0
2 # "covse13" <- 0
3 # "covse23" <- 0
4 # "covsp12" <- 0
5 # "covsp13" <- 0
6 # "covsp23" <- 0
```

So that they look like:

```
1 "covse12" <- 0
2 # "covse13" <- 0
3 # "covse23" <- 0
4 "covsp12" <- 0
5 # "covsp13" <- 0
6 # "covsp23" <- 0
```

```
1 results <- run.jags('auto3tihw.bug')
```

## Exercise

Simulate data with  $N=1000$  and dependence between tests 1 and 2

Then fit a model assuming independence between all tests and compare the results to your simulation parameters

Now turn on covariance between tests 1 and 2 and refit the model.  
Are the results more reasonable?

## Optional Exercise

Re-fit a model to this data using all three possible  $\text{covse}$  and  $\text{covsp}$  parameters

What do you notice about the results?

## **Session 3c: Model selection**

---

# Motivation

- Choosing between candidate models
  - DIC
  - Bayes Factors
  - BIC
  - WAIC
  - Effect size spans zero?

# Motivation

- Choosing between candidate models
  - DIC
  - Bayes Factors
  - BIC
  - WAIC
  - Effect size spans zero?
- Assessing model adequacy:
  - Verify using a simulation study
  - Posterior predictive p-values
  - Comparison of results from different models eg:
    - Independence vs covariance
    - Different priors



# Motivation

- Choosing between candidate models
  - DIC
  - Bayes Factors
  - BIC
  - WAIC
  - Effect size spans zero?
- Assessing model adequacy:
  - Verify using a simulation study
  - Posterior predictive p-values
  - Comparison of results from different models eg:
    - Independence vs covariance
    - Different priors

Others?

- DIC
  - Works well for hierarchical normal models
  - To calculate:
    - Add dic and ped to the monitors in runjags
    - But be cautious with these types of models

# DIC and WAIC

- DIC
  - Works well for hierarchical normal models
  - To calculate:
    - Add dic and ped to the monitors in runjags
    - But be cautious with these types of models
- WAIC
  - Approximation to LOO
  - Needs independent likelihoods
    - Could work for individual-level models?
  - Currently a pain to calculate
    - See WAIC.R in the GitHub directory
    - And/or wait for updates to runjags (and particularly JAGS 5)

# DIC and WAIC

- DIC
  - Works well for hierarchical normal models
  - To calculate:
    - Add dic and ped to the monitors in runjags
    - But be cautious with these types of models
- WAIC
  - Approximation to LOO
  - Needs independent likelihoods
    - Could work for individual-level models?
  - Currently a pain to calculate
    - See WAIC.R in the GitHub directory
    - And/or wait for updates to runjags (and particularly JAGS 5)

```
1 install.packages('runjags',  
  ↪ repos=c("https://ku-awdc.github.io/drat/",  
  ↪ "https://cran.rstudio.com/"))
```

## Some advice

- Always start by simulating data and verifying that you can recover the parameters
  - The simulation can be more complex than the model!
  - See the `autorun.jags` function
- If you have different candidate models then compare the posteriors between models

## Some advice

- Always start by simulating data and verifying that you can recover the parameters
  - The simulation can be more complex than the model!
  - See the `autorun.jags` function
- If you have different candidate models then compare the posteriors between models
- A particular issue is test dependence
  - Is there biological justification for the correlation?
  - Are the test sensitivity/specificity estimates consistent?
  - Do the `covse` / `covsp` estimates overlap zero?

## Some advice

- Always start by simulating data and verifying that you can recover the parameters
  - The simulation can be more complex than the model!
  - See the `autorun.jags` function
- If you have different candidate models then compare the posteriors between models
- A particular issue is test dependence
  - Is there biological justification for the correlation?
  - Are the test sensitivity/specificity estimates consistent?
  - Do the `covse` / `covsp` estimates overlap zero?
- Any other good advice?!?

## Free practical time

- Explore the optional exercises (and solutions) and feel free to ask questions!
- Feedback very welcome!