

# Hands-on training session 3

Hui-Walter models with more than two diagnostic tests

---

Matt Denwood   Giles Innocent   Sonja Hartnack

2020-02-19

# Introduction

---

# Overview

Date/time:

- 20th February 2020
- 14.00 - 15.30

Teachers:

- Matt Denwood (presenter)
- Giles Innocent
- Sonja Hartnack

# Recap

- JAGS / runjags is the easy way to work with complex models
  - But we *still have to* check convergence and effective sample size!
- Estimating sensitivity and specificity is like pulling a rabbit out of a hat
  - Multiple populations helps **a lot**
  - Strong priors for one of the tests helps even more!

# Recap

- JAGS / runjags is the easy way to work with complex models
  - But we *still have to* check convergence and effective sample size!
- Estimating sensitivity and specificity is like pulling a rabbit out of a hat
  - Multiple populations helps **a lot**
  - Strong priors for one of the tests helps even more!
- But what if the tests are not independent of each other?

## **Session 3a: Hui-Walter models for multiple conditionally independent tests**

---

# What exactly is our latent class?

- What do we mean by “conditionally independent?”

# What exactly is our latent class?

- What do we mean by “conditionally independent?”
- Example: we have three antibody tests
  - The latent status is actually ‘producing antibodies’ not ‘diseased’



# What exactly is our latent class?

- What do we mean by “conditionally independent?”
- Example: we have three antibody tests
  - The latent status is actually ‘producing antibodies’ not ‘diseased’
- Example: antibody vs egg count tests for liver fluke
  - Does the latent state include migrating juvenile fluke?

# What exactly is our latent class?

- What do we mean by “conditionally independent?”
- Example: we have three antibody tests
  - The latent status is actually ‘producing antibodies’ not ‘diseased’
- Example: antibody vs egg count tests for liver fluke
  - Does the latent state include migrating juvenile fluke?
- We’re actually pulling **something** out of a hat, and deciding to call it a rabbit

# Simulating data

Simulating data using an arbitrary number of independent tests is quite straightforward.

```
1  # Parameter values to simulate:
2  N <- 200
3  se1 <- 0.8
4  sp1 <- 0.95
5  se2 <- 0.9
6  sp2 <- 0.99
7  se3 <- 0.95
8  sp3 <- 0.95
9
10 Populations <- 2
11 prevalence <- c(0.25,0.75)
12 Group <- sample(1:Populations, N, replace=TRUE)
```

```

1  # Ensure replicable data:
2  set.seed(2020-02-18)
3
4  # Simulate the true latent state (which is unobserved in real
   ↪ life):
5  true <- rbinom(N, 1, prevalence[Group])
6  # Simulate test results for test 1:
7  test1 <- rbinom(N, 1, se1*true + (1-sp1)*(1-true))
8  # Simulate test results for test 2:
9  test2 <- rbinom(N, 1, se2*true + (1-sp2)*(1-true))
10 # Simulate test results for test 3:
11 test3 <- rbinom(N, 1, se3*true + (1-sp3)*(1-true))
12
13 simdata <- data.frame(Population=factor(Group), Test1=test1,
   ↪ Test2=test2, Test3=test3)

```

# Model specification

- Like for two tests, except it is now a 2x2x2 table
  - If calculating this manually, take **extreme** care with multinomial tabulation
- Or use `autohuiwalter`
  - This will also deal gracefully with missing data in one or more test results

```
1 source("autohuiwalter.R")
2 auto_huiwalter(simdata[,c('Population', 'Test1', 'Test2', 'Test3')],
  ↪ outfile='auto3thw.bug')
```

```

1  for(p in 1:Populations){
2  Tally_RRR[1:8,p] ~ dmulti(prob_RRR[1:8,p], N_RRR[p])
3  prob_RRR[1:8,p] <- se_prob[1:8,p] + sp_prob[1:8,p]
4
5  # Probability of observing Test1- Test2- Test3- from a true
   ↪ positive::
6  se_prob[1,p] <- prev[p] * ((1-se[1])*(1-se[2])*(1-se[3])) +covse12
   ↪ +covse13 +covse23)
7  # Probability of observing Test1- Test2- Test3- from a true
   ↪ negative::
8  sp_prob[1,p] <- (1-prev[p]) * (sp[1]*sp[2]*sp[3] +covsp12
   ↪ +covsp13 +covsp23)
9
10 . . .
11
12 # Probability of observing Test1+ Test2+ Test3+ from a true
   ↪ positive::
13 se_prob[8,p] <- prev[p] * (se[1]*se[2]*se[3] +covse12 +covse13
   ↪ +covse23)
14 # Probability of observing Test1+ Test2+ Test3+ from a true
   ↪ negative::
15 sp_prob[8,p] <- (1-prev[p]) * ((1-sp[1])*(1-sp[2])*(1-sp[3]))
   ↪ +covsp12 +covsp13 +covsp23)
16 }

```

## Alternative model specification

We might want to explicitly model the latent state:

```
1  model{  
2  
3    for(i in 1:N){  
4      truestatus[i] ~ dbern(prev[Population[i]])  
5  
6      Status[i] ~ dcat(prob[i, ])  
7      prob[1:8,i] <- se_prob[1:8,i] + sp_prob[1:8,i]
```

```

1
2     se_prob[1,p] <- truestatus[i] *
   ↪   ((1-se[1])*(1-se[2])*(1-se[3]))
3     sp_prob[1,p] <- (1-truestatus[i]) * (sp[1]*sp[2]*sp[3])
4
5     se_prob[2,p] <- truestatus[i] *
   ↪   (se[1]*(1-se[2])*(1-se[3]))
6     sp_prob[2,p] <- (1-truestatus[i]) *
   ↪   ((1-sp[1])*sp[2]*sp[3])
7
8     se_prob[3,p] <- truestatus[i] *
   ↪   ((1-se[1])*se[2]*(1-se[3]))
9     sp_prob[3,p] <- (1-truestatus[i]) *
   ↪   (sp[1]*(1-sp[2])*sp[3])
10
11     se_prob[4,p] <- truestatus[i] * (se[1]*se[2]*(1-se[3]))
12     sp_prob[4,p] <- (1-truestatus[i]) *
   ↪   ((1-sp[1])*(1-sp[2])*sp[3])
13
14     se_prob[5,p] <- truestatus[i] *
   ↪   ((1-se[1])*(1-se[2])*se[3])
15     sp_prob[5,p] <- (1-truestatus[i]) *
   ↪   (sp[1]*sp[2]*(1-sp[3]))

```



```

1
2     prev[1] ~ dbeta(1,1)
3     prev[2] ~ dbeta(1,1)
4
5     se[1] ~ dbeta(1, 1)T(1-sp[1], )
6     sp[1] ~ dbeta(1, 1)
7     se[2] ~ dbeta(1, 1)T(1-sp[2], )
8     sp[2] ~ dbeta(1, 1)
9     se[3] ~ dbeta(1, 1)T(1-sp[3], )
10    sp[3] ~ dbeta(1, 1)
11
12    #data# Status, N, Population
13    #monitor# prev, se, sp
14    #inits# prev, se, sp
15    }

```

But this is inefficient

There is also no way to distinguish individuals within the same boxes

We could also use the estimated se/sp/prev to post-calculate these status probabilities

This is useful for post-hoc ROC

## Exercise

Simulate data from 3 tests and analyse using the `autohuiwalter` function

Do the estimates of  $Se/Sp$  correspond to the simulation parameters?

Make some data missing for one or more tests and re-generate the model

- Can you see what has changed in the code?

## Optional Exercise

Simulate data from 3 antibody tests with ab positive step [give code in solution]

Does the se/sp estimated by a model recover the parameters?

Why not?

## **Session 3b: Hui-Walter models for multiple tests with conditional dependence**

---

## Branching of processes leading to test results

Example: two antibody tests and one antigen test

Or three antibody tests where one has a different target to the other two

# Simulating data

It helps to consider the data simulation as a biological process.

```
1  # Parameter values to simulate:
2  N <- 200
3  se1 <- 0.8; sp1 <- 0.95
4  se2 <- 0.9; sp2 <- 0.99
5  se3 <- 0.95; sp3 <- 0.95
6
7  Populations <- 2
8  prevalence <- c(0.25,0.75)
9  Group <- rep(1:Populations, each=N)
10
11 # Ensure replicable data:
12 set.seed(2017-11-21)
13
14 # The probability of an antibody response given disease:
15 abse <- 0.8
16 # The probability of no antibody response given no disease:
17 absp <- 1 - 0.2
```

```

1  # Simulate the true latent state:
2  true <- rbinom(N*Populations, 1, prevalence[Group])
3  # Simulate test results for test 1:
4  test1 <- rbinom(N*Populations, 1, se1*true + (1-sp1)*(1-true))
5  # Tests 2 & 3 will be co-dependent on antibody response:
6  antibody <- rbinom(N*Populations, 1, abse*true +
   ↪ (1-absp)*(1-true))
7  # Simulate test 2 & 3 results based on this other latent state:
8  test2 <- rbinom(N*Populations, 1, se2*antibody +
   ↪ (1-sp2)*(1-antibody))
9  test3 <- rbinom(N*Populations, 1, se3*antibody +
   ↪ (1-sp3)*(1-antibody))
10
11 ind3tests <- data.frame(Population=Group, Test1=test1,
   ↪ Test2=test2, Test3=test3)

```



```
1  # The overall sensitivity of the correlated tests is:
2  abse*se2 + (1-abse)*(1-sp2)
```

```
1  ## [1] 0.722
```

```
1  abse*se3 + (1-abse)*(1-sp3)
```

```
1  ## [1] 0.77
```

```
1  # The overall specificity of the correlated tests is:
2  absp*sp2 + (1-absp)*(1-se2)
```

```
1  ## [1] 0.812
```

```
1  absp*sp3 + (1-absp)*(1-se3)
```

```
1  ## [1] 0.77
```

```
1  # The overall sensitivity of the correlated tests is:
2  abse*se2 + (1-abse)*(1-sp2)
```

```
1  ## [1] 0.722
```

```
1  abse*se3 + (1-abse)*(1-sp3)
```

```
1  ## [1] 0.77
```

```
1  # The overall specificity of the correlated tests is:
2  absp*sp2 + (1-absp)*(1-se2)
```

```
1  ## [1] 0.812
```

```
1  absp*sp3 + (1-absp)*(1-se3)
```

```
1  ## [1] 0.77
```

We need to think carefully about what we are conditioning on when interpreting sensitivity and specificity!

# Model specification

```
1      se_prob[1,p] <- prev[p] * ((1-se[1])*(1-se[2])*(1-se[3])
↪      +covse12 +covse13 +covse23)
2      sp_prob[1,p] <- (1-prev[p]) * (sp[1]*sp[2]*sp[3] +covsp12
↪      +covsp13 +covsp23)
3
4      se_prob[2,p] <- prev[p] * (se[1]*(1-se[2])*(1-se[3]) -covse12
↪      -covse13 +covse23)
5      sp_prob[2,p] <- (1-prev[p]) * ((1-sp[1])*sp[2]*sp[3] -covsp12
↪      -covsp13 +covsp23)
6
7      ...
8
9      # Covariance in sensitivity between tests 1 and 2:
10     covse12 ~ dunif( (se[1]-1)*(1-se[2]) , min(se[1],se[2]) -
↪     se[1]*se[2] )
11
12     # Covariance in specificity between tests 1 and 2:
13     covsp12 ~ dunif( (sp[1]-1)*(1-sp[2]) , min(sp[1],sp[2]) -
↪     sp[1]*sp[2] )
14
15     ...
```

# Generating the model

Use `autohuiwalter` with argument `covon=TRUE`

```
1 source('autohuiwalter.R')
2 auto_huiwalter(ind3tests, 'auto3tihw.bug', covon=TRUE)

1 # Covariance in sensitivity between Test1 and Test2 tests:
2 covse12 ~ dunif( (se[1]-1)*(1-se[2]) , min(se[1],se[2]) -
  ↪ se[1]*se[2] ) ## if the sensitivity of these tests may be
  ↪ correlated
3 # covse12 <- 0 ## if the sensitivity of these tests can be
  ↪ assumed to be independent
4 # Covariance in specificity between Test1 and Test2 tests:
5 covsp12 ~ dunif( (sp[1]-1)*(1-sp[2]) , min(sp[1],sp[2]) -
  ↪ sp[1]*sp[2] ) ## if the specificity of these tests may be
  ↪ correlated
6 # covsp12 <- 0 ## if the specificity of these tests can be
  ↪ assumed to be independent
```

```

1  # Covariance in sensitivity between Test1 and Test3 tests:
2  covse13 ~ dunif( (se[1]-1)*(1-se[3]) , min(se[1],se[3]) -
   ↪ se[1]*se[3] ) ## if the sensitivity of these tests may be
   ↪ correlated
3  # covse13 <- 0 ## if the sensitivity of these tests can be
   ↪ assumed to be independent
4  # Covariance in specificity between Test1 and Test3 tests:
5  covsp13 ~ dunif( (sp[1]-1)*(1-sp[3]) , min(sp[1],sp[3]) -
   ↪ sp[1]*sp[3] ) ## if the specificity of these tests may be
   ↪ correlated
6  # covsp13 <- 0 ## if the specificity of these tests can be
   ↪ assumed to be independent
7
8  # Covariance in sensitivity between Test2 and Test3 tests:
9  covse23 ~ dunif( (se[2]-1)*(1-se[3]) , min(se[2],se[3]) -
   ↪ se[2]*se[3] ) ## if the sensitivity of these tests may be
   ↪ correlated
10 # covse23 <- 0 ## if the sensitivity of these tests can be
    ↪ assumed to be independent
11 # Covariance in specificity between Test2 and Test3 tests:
12 covsp23 ~ dunif( (sp[2]-1)*(1-sp[3]) , min(sp[2],sp[3]) -
   ↪ sp[2]*sp[3] ) ## if the specificity of these tests may be
   ↪ correlated
13 # covsp23 <- 0 ## if the specificity of these tests can be

```

## Exercise

Simulate data with a dependence between 2 tests

Model assuming conditional independence biases the estimates

Turn on covariance between the two tests

Model with conditional dependence has bigger CI but unbiased

## Optional Exercise

Activate covariance between all 3 tests

## **Session 3c: Model selection**

---



# Motivation

- Choosing between candidate models
  - DIC
  - Bayes Factors
  - BIC
  - WAIC

# Motivation

- Choosing between candidate models
  - DIC
  - Bayes Factors
  - BIC
  - WAIC
- Assessing model adequacy:
  - Verify using a simulation study
  - Posterior predictive p-values
  - Comparison of results from different models eg:
    - Independence vs covariance
    - Different priors

# Motivation

- Choosing between candidate models
  - DIC
  - Bayes Factors
  - BIC
  - WAIC
- Assessing model adequacy:
  - Verify using a simulation study
  - Posterior predictive p-values
  - Comparison of results from different models eg:
    - Independence vs covariance
    - Different priors

Others? Discussion!

# DIC and WAIC

- DIC
  - Works well for hierarchical normal models
  - To calculate:
    - Add dic and ped to the monitors in runjags
    - But don't trust the results for these types of models
- WAIC
  - Approximation to LOO
  - Needs independent likelihoods
    - Could work for individual-level models?
  - Currently a pain to calculate
    - See WAIC.R in the GitHub directory
    - And/or wait for updates to runjags (and particularly JAGS 5)

# DIC and WAIC

- DIC
  - Works well for hierarchical normal models
  - To calculate:
    - Add dic and ped to the monitors in runjags
    - But don't trust the results for these types of models
- WAIC
  - Approximation to LOO
  - Needs independent likelihoods
    - Could work for individual-level models?
  - Currently a pain to calculate
    - See WAIC.R in the GitHub directory
    - And/or wait for updates to runjags (and particularly JAGS 5)

```
1 install.packages('runjags',  
  ↪ repos=c("https://ku-awdc.github.io/drat/",  
  ↪ "https://cran.rstudio.com/"))
```

## Discussion and free practical time

Any questions?